

```
In [1]: import pandas as pd
```

```
In [4]: import numpy as np
```

```
In [11]: import matplotlib.pyplot as plt
```

```
In [12]: import seaborn as sb
```

```
In [13]: sb.set(color_codes= True)
```

```
In [16]: import warnings
```

```
In [20]: warnings.filterwarnings('ignore')
```

Importing Datasets

```
In [23]: client_data = pd.read_csv('client_data.csv')
client_data .head()
```

```
Out[23]:
```

	id	channel_sales	cons_12m	cons_gas_12m
0	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcsosbicdxkicaua	0	5494€
1	d29c2c54acc38ff3c0614d0a653813dd	MISSING	4660	C
2	764c75f661154dac3a6c254cd082ea7d	foosdfpfkusacimwkcsosbicdxkicaua	544	C
3	bba03439a292a1e166f80264c16191cb	lmkebamcaclubfxadlmueccxoimlema	1584	C
4	149d57cf92fc41cf94415803a877cb4b	MISSING	4425	C

5 rows × 26 columns

```
In [24]: price_data = pd.read_csv('price_data.csv')
price_data .head()
```

Out[24]:

	id	price_date	price_off_peak_var	price_peak_var	price_mid_peak
0	038af19179925da21a25619c5a24b745	01-01-2015	0.151367	0.0	
1	038af19179925da21a25619c5a24b745	01-02-2015	0.151367	0.0	
2	038af19179925da21a25619c5a24b745	01-03-2015	0.151367	0.0	
3	038af19179925da21a25619c5a24b745	01-04-2015	0.149626	0.0	
4	038af19179925da21a25619c5a24b745	01-05-2015	0.149626	0.0	

3: Descriptive Statistics of Data

Data Types

```
In [28]: print("rows and column are ",client_data.shape)
rows and column are (14606, 26)
```

```
In [29]: print("rows and column are ",price_data.shape)
rows and column are (193002, 8)
```

Basic info of client_data

```
In [30]: client_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14606 entries, 0 to 14605
Data columns (total 26 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   id                                         14606 non-null  object
1   channel_sales                             14606 non-null  object
2   cons_12m                                  14606 non-null  int64
3   cons_gas_12m                              14606 non-null  int64
4   cons_last_month                           14606 non-null  int64
5   date_activ                                14606 non-null  object
6   date_end                                  14606 non-null  object
7   date_modif_prod                           14606 non-null  object
8   date_renewal                              14606 non-null  object
9   forecast_cons_12m                         14606 non-null  float64
10  forecast_cons_year                         14606 non-null  int64
11  forecast_discount_energy                   14606 non-null  int64
12  forecast_meter_rent_12m                   14606 non-null  float64
13  forecast_price_energy_off_peak             14606 non-null  float64
14  forecast_price_energy_peak                 14606 non-null  float64
15  forecast_price_pow_off_peak                14606 non-null  float64
16  has_gas                                    14606 non-null  object
17  imp_cons                                   14606 non-null  float64
18  margin_gross_pow_ele                       14606 non-null  float64
19  margin_net_pow_ele                         14606 non-null  float64
20  nb_prod_act                               14606 non-null  int64
21  net_margin                                14606 non-null  float64
22  num_years_antig                            14606 non-null  int64
23  origin_up                                 14606 non-null  object
24  pow_max                                    14606 non-null  float64
25  churn                                      14606 non-null  int64
dtypes: float64(10), int64(8), object(8)
memory usage: 2.9+ MB

```

Basic info of price_data


In [32]: `price_data.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 193002 entries, 0 to 193001
Data columns (total 8 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   id                                         193002 non-null  object
1   price_date                                193002 non-null  object
2   price_off_peak_var                        193002 non-null  float64
3   price_peak_var                            193002 non-null  float64
4   price_mid_peak_var                        193002 non-null  float64
5   price_off_peak_fix                        193002 non-null  float64
6   price_peak_fix                            193002 non-null  float64
7   price_mid_peak_fix                        193002 non-null  float64
dtypes: float64(6), object(2)
memory usage: 11.8+ MB

```

In [33]: `sb.heatmap(client_data.isnull())`
`mpl.show()`




```
In [34]: client_data.isnull().sum()
```

```
Out[34]: id 0
channel_sales 0
cons_12m 0
cons_gas_12m 0
cons_last_month 0
date_activ 0
date_end 0
date_modif_prod 0
date_renewal 0
forecast_cons_12m 0
forecast_cons_year 0
forecast_discount_energy 0
forecast_meter_rent_12m 0
forecast_price_energy_off_peak 0
forecast_price_energy_peak 0
forecast_price_pow_off_peak 0
has_gas 0
imp_cons 0
margin_gross_pow_ele 0
margin_net_pow_ele 0
nb_prod_act 0
net_margin 0
num_years_antig 0
origin_up 0
pow_max 0
churn 0
dtype: int64
```

checking missing values of price data

```
In [35]: sb.heatmap(price_data.isnull())
mpl.show()
```



```
In [36]: price_data .isnull().sum()
```

```
Out[36]: id                0
price_date              0
price_off_peak_var      0
price_peak_var          0
price_mid_peak_var      0
price_off_peak_fix      0
price_peak_fix          0
price_mid_peak_fix      0
dtype: int64
```

As we can see there is no missing value there in both the tables

Statistics

Now let see some statistics about both the data tables

```
In [40]: client_data.describe().T
```

Out[40]:

	count	mean	std	min	25%	50%	75%	max
cons_12m	14606.0	159220.286252	573465.264198	0.0	5674.750000	14115.500		
cons_gas_12m	14606.0	28092.375325	162973.059057	0.0	0.000000	0.000		
cons_last_month	14606.0	16090.269752	64364.196422	0.0	0.000000	792.500		
forecast_cons_12m	14606.0	1868.614880	2387.571531	0.0	494.995000	1112.875		
forecast_cons_year	14606.0	1399.762906	3247.786255	0.0	0.000000	314.000		
forecast_discount_energy	14606.0	0.966726	5.108289	0.0	0.000000	0.000		
forecast_meter_rent_12m	14606.0	63.086871	66.165783	0.0	16.180000	18.795		
forecast_price_energy_off_peak	14606.0	0.137283	0.024623	0.0	0.116340	0.143		
forecast_price_energy_peak	14606.0	0.050491	0.049037	0.0	0.000000	0.084		
forecast_price_pow_off_peak	14606.0	43.130056	4.485988	0.0	40.606701	44.311		
imp_cons	14606.0	152.786896	341.369366	0.0	0.000000	37.395		
margin_gross_pow_ele	14606.0	24.565121	20.231172	0.0	14.280000	21.640		
margin_net_pow_ele	14606.0	24.562517	20.230280	0.0	14.280000	21.640		
nb_prod_act	14606.0	1.292346	0.709774	1.0	1.000000	1.000		
net_margin	14606.0	189.264522	311.798130	0.0	50.712500	112.530		
num_years_antig	14606.0	4.997809	1.611749	1.0	4.000000	5.000		
pow_max	14606.0	18.135136	13.534743	3.3	12.500000	13.856		
churn	14606.0	0.097152	0.296175	0.0	0.000000	0.000		

1) client_data.describe() has statistics as rows and features as columns.

2) client_data.describe().T has features as rows and statistics as columns.

```
In [42]: # statistical summary of price data
price_data.describe().T
```

Out[42]:

	count	mean	std	min	25%	50%	75%	max
price_off_peak_var	193002.0	0.141027	0.025032	0.0	0.125976	0.146033	0.151635	0.28070
price_peak_var	193002.0	0.054630	0.049924	0.0	0.000000	0.085483	0.101673	0.22978
price_mid_peak_var	193002.0	0.030496	0.036298	0.0	0.000000	0.000000	0.072558	0.11410
price_off_peak_fix	193002.0	43.334477	5.410297	0.0	40.728885	44.266930	44.444710	59.44471
price_peak_fix	193002.0	10.622875	12.841895	0.0	0.000000	0.000000	24.339581	36.49069
price_mid_peak_fix	193002.0	6.409984	7.773592	0.0	0.000000	0.000000	16.226389	17.45822

Data visualization

```
In [49]: def plot_stacked_bars(dataframe, title_, size_=(18, 10), rot_=0, legend_="upper right")
        """
```

```

Plot stacked bars with annotations
"""
ax = dataframe.plot(
    kind="bar",
    stacked=True,
    figsize=size_,
    rot=rot_,
    title=title_
)

# Annotate bars
annotate_stackedBars(ax, textsize=14)
# Rename Legend
plt.legend(["Retention", "Churn"], loc=legend_)
# Labels
plt.ylabel("Company base (%)")
plt.show()

def annotate_stackedBars(ax, pad=0.99, colour="white", textsize=13):
    """
    Add value annotations to the bars
    """

    # Iterate over the plotted rectangles/bars
    for p in ax.patches:

        # Calculate annotation
        value = str(round(p.get_height(),1))
        # If value is 0 do not annotate
        if value == '0.0':
            continue
        ax.annotate(
            value,
            ((p.get_x()+ p.get_width()/2)*pad-0.05, (p.get_y()+p.get_height()/2)*pa
            color=colour,
            size=textsize
        )

```

```

In [50]: churn =client_data[['id','churn']]
churn.column =['companies','churn']
churn_total = churn.groupby(churn['churn']).count()
churn_percentage = churn_total/churn_total.sum()*100

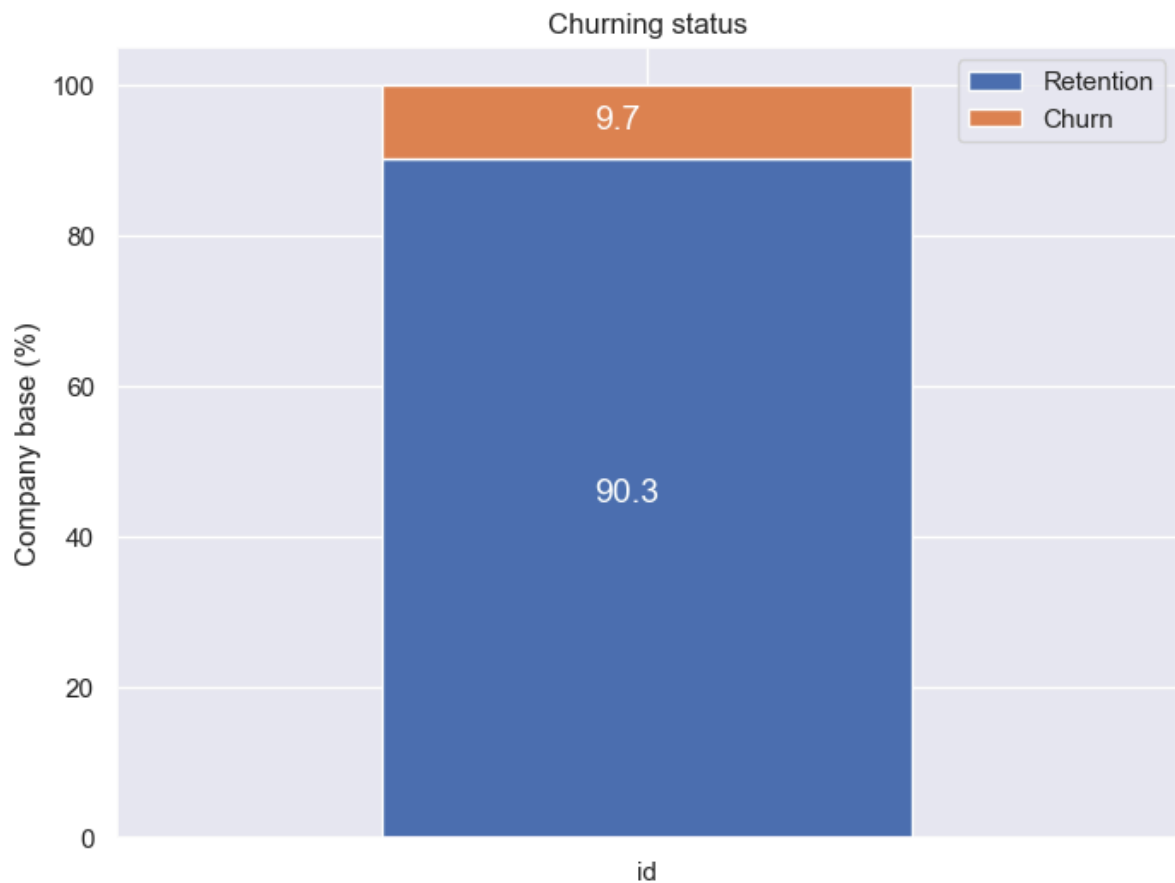
```

```

In [51]: plot_stackedBars(churn_percentage.transpose(), "Churning status", (8, 6), legend_=

print("\n ----- Value Counts ----- \n")
print(client_data['churn'].value_counts())

```


----- Value Counts -----

```
0    13187
1     1419
Name: churn, dtype: int64
```

Insight: nearly 10% of users have churned

```
In [59]: channel = client_data[['id', 'channel_sales', 'churn']]
channel = channel.groupby([channel['channel_sales'], channel['churn']])['id'].count
channel_churn = (channel.div(channel.sum(axis=1), axis=0) * 100).sort_values(by=[1])
```

```
In [60]: plot_stacked_bars(channel_churn, 'Sales channel', rot_=30)
```



Consumption

Let's see the distribution of the consumption in the last year and month. Since the consumption data is univariate, let's use histograms to visualize their distribution.

In [61]: `consumption = client_data[['id', 'cons_12m', 'cons_gas_12m', 'cons_last_month', 'imp_cons', 'has_churn']]`
`consumption.head()`

Out[61]:

	id	cons_12m	cons_gas_12m	cons_last_month	imp_cons	has_churn
0	24011ae4ebbe3035111d65fa7c15bc57	0	54946	0	0.00	
1	d29c2c54acc38ff3c0614d0a653813dd	4660	0	0	0.00	
2	764c75f661154dac3a6c254cd082ea7d	544	0	0	0.00	
3	bba03439a292a1e166f80264c16191cb	1584	0	0	0.00	
4	149d57cf92fc41cf94415803a877cb4b	4425	0	526	52.32	

In [65]: `def plot_distribution(dataframe, column, ax, bins_=70):`
`# Create a temporal dataframe with the data to be plot`
`temp = pd.DataFrame({"Retention": dataframe[dataframe["churn"]==0][column],`
`"Churn": dataframe[dataframe["churn"]==1][column]})`
`# Plot the histogram`
`temp[["Retention", "Churn"]].plot(kind='hist', bins=bins_, ax=ax, stacked=True)`
`# X-axis label`
`ax.set_xlabel(column)`
`# Change the x-axis to plain style`
`ax.ticklabel_format(style='plain', axis='x')`

In [67]: `fig, axes = plt.subplots(nrows=4, figsize=(18, 25))`
`plot_distribution(consumption, 'imp_cons', axes[3])`

```

plot_distribution(consumption, 'cons_12m', axs[0])
plot_distribution(consumption, 'cons_last_month', axs[2])
plot_distribution(consumption[consumption['has_gas'] == 't'], 'cons_gas_12m', axs[1])

mpl.savefig("consumption Distribution.png", bbox_inches="tight")

```



Forecast

```

In [68]: forecast = client_data[["id", "forecast_cons_12m", "forecast_cons_year", "forecast_c",
    "forecast_price_energy_off_peak", "forecast_price_energy_peak", "forecast_price_

forecast.head()

```

Out[68]:


	id	forecast_cons_12m	forecast_cons_year	forecast_discount_ene
0	24011ae4ebbe3035111d65fa7c15bc57	0.00	0	
1	d29c2c54acc38ff3c0614d0a653813dd	189.95	0	
2	764c75f661154dac3a6c254cd082ea7d	47.96	0	
3	bba03439a292a1e166f80264c16191cb	240.04	0	
4	149d57cf92fc41cf94415803a877cb4b	445.75	526	

In [70]:

```
fig, axs = plt.subplots(nrows=7, figsize=(35,50))

# Plot histogram
plot_distribution(client_data, "forecast_cons_12m", axs[6])
plot_distribution(client_data, "forecast_cons_year", axs[5])
plot_distribution(client_data, "forecast_discount_energy", axs[4])
plot_distribution(client_data, "forecast_meter_rent_12m", axs[3])
plot_distribution(client_data, "forecast_price_energy_off_peak", axs[2])
plot_distribution(client_data, "forecast_price_energy_peak", axs[1])
plot_distribution(client_data, "forecast_price_pow_off_peak", axs[0])


plt.savefig("Forecast Views .png", bbox_inches="tight")
```



Contract type

```
In [71]: contract_type = client_data[['id', 'has_gas', 'churn']]
contract = contract_type.groupby([contract_type['churn'], contract_type['has_gas']])
contract_percentage = (contract.div(contract.sum(axis=1), axis=0) * 100).sort_values
```

```
In [73]: plot_stackedBars(contract_percentage, 'Contract type with Gas')
```




Subscribed power

```
In [74]: power = client_data[['id', 'pow_max', 'churn']]
```

```
In [79]: fig, axs = plt.subplots(nrows=1, figsize=(20, 18))
plot_distribution(power, 'pow_max', axs)

plt.savefig("Subscribed power view .png", bbox_inches="tight")
```



In []: